

START OF QUIZ

Student ID:

38419826,Zeng,Lingsong

Question 1

Topic: Lecture 8

Source: Lecture 8

Where do you think pragmatic learning (ie, intent) might fall within the layers of an LLM?
Explain briefly. How might we test for it? (1)

Question 2

Topic: Lecture 6

Source: Lecture 6

Explain how an MSD differs from a POS tag, and how it's similar. (1)

Question 3

Topic: Lecture 7

Source: Lecture 7

When generating inflections, what features of the input do you think the model most attends to? (1)

Question 4

Topic: Lecture 6

Source: Lecture 6

Would F1 score be an appropriate measure for gauging the quality of a morphological analyzer? Explain why or why not. (1)

Question 5

Topic: Lecture 5

Source: Lecture 5

Imagine that we wanted to create POS embeddings in the same way that we create word embeddings. Explain one way that it would be very similar to word embeddings, and one that it would be very different. (1)

Question 6

Topic: Lecture 8

Source: Lecture 8

Imagine we were designing a probe to understand whether a model were gender biased. How might we design such a probe, and if we found the model to exhibit such a bias, what suggestions would you make to neutralize the bias? (2)

Question 7

Topic: Lecture 5

Source: Lecture 5

Imagine that we have some pre-trained multilingual embeddings of really high quality. We train a POS tagger for a very common language, with lots of data, embedding the data with the multilingual embeddings. At inference, we then replace the input with another language. Do you think the tagger would beat a majority baseline? Explain your reasoning, and list any assumptions. (2)

Question 8

Topic: Lecture 7

Source: Lecture 7

What benefits might encoding MSDs with a second encoder have over a single encoder approach? Can you think of any disadvantages? (2)

Question 9

Topic: Long

Source: Lecture 5

In class, we talked about how POS and morphological information is often latently encoded in word embeddings, but not in character embeddings. Let's think about subword embeddings, since most DL models are going to use subword representations. If a word is split, where do you think this information is encoded, and does it matter? Explain your reasoning.
(3)

END OF QUIZ