

START OF QUIZ

Student ID:

67446104,Nguyen,Minh

Question 1

Topic: Lecture 7

Source: Lecture 7

We took a look at 2 different ways of implementing the TextTiling algorithm - one with vector overlap, and one with BERT. Can you think of how we might modify the algorithm further to strengthen up its weaknesses? (No is not a valid answer.) (2)

Question 2

Topic: Lecture 6

Source: Lecture 6

When running a window-based approach to vector embeddings (such as CBOW or skip-gram), when would it make sense to keep stopwords, and when would it make sense to remove them? (1)

Question 3

Topic: Lecture 8

Source: Lecture 8

Why is it easy to create negative examples for lexical coherence tests? (1)

Question 4

Topic: Lecture 5

Source: Lecture 5

Which is likely to have the lowest PMI? A rare word and a frequent word that appear together frequently, or two frequent words that appear together frequently? (1)

Question 5

Topic: Lecture 6

Source: Lecture 6

What is the purpose of negative sampling in a Word2Vec model? (1)

Question 6

Topic: Lecture 8

Source: Lecture 8

Do you think we could use word embeddings for coreference resolution? What kind of assumptions would we be making, and why do you think it might still be a very difficult task?
(2)

Question 7

Topic: Lecture 7

Source: Lecture 7

How is the TextTiling algorithm similar to the Lesk algorithm? How is it different? (2)

Question 8

Topic: Lecture 5

Source: Lecture 5

What is PMI measuring? That is, what does it mean for two words to have high PMI? (1)

Question 9

Topic: Long

Source: Lecture 5

All of these embeddings we've been looking at have been an effort to translate meaning into math, so that we can use computational algorithms (which are good at math) to process meaning. To what extent do you think that these are a good approximation for how we understand language, and to what extent do you think they are a poor approximation? (3)

END OF QUIZ