START OF QUIZ
Student ID:
90412503,Yin,Ting

Question 1

Topic: Lecture 5
Source: Lecture 5

ChatGPT differs significantly from even other neural Q/A systems. Provide at least 2 significant differences, and briefly describe them. (2)

Question 2

Topic: Lecture 5
Source: Lecture 5

Describe the two ways that we can construct Q/A databases, and how they differ. (2)

Question 3

Topic: Lecture 8
Source: Lecture 8

How is it that Eliza can use words / phrases that she doesn't have in her templates? (1)

Question 4

Topic: Lecture 7
Source: Lecture 7

We discussed slot error rate in class, but it's fully-supervised. Can you think of a distantly-supervised way to calculate essentially the same thing? (1)

Question 5

Topic: Lecture 6
Source: Lecture 6

What tools are necessary to extract an RDF triple from a question? Provide at least 2, and briefly explain. (1)

Question 6

Topic: Lecture 8
Source: Lecture 8

Do you think a dialogue policy state graph is a Markov Chain? Briefly describe why or why not. (If you can't remember Markov chains, we talked about them in DSCI 572). (1)

Question 7

Topic: Lecture 7
Source: Lecture 7

Generate a frame for a “recommend a movie” dialogue action. It should have at least 5 slots to fill. (2)

Question 8

Topic: Lecture 6
Source: Lecture 6

Provide a reasonable logical representation of the question "Who starred in Casablanca?" (1)

Question 9

Topic: Long
Source: Lecture 7

Imagine that I'm working with a client who wants a dialogue system that provides advice for his company. It has to fit on a phone, but might end up in regions with very limited cell service, so it has to be locally installed. We have limited memory (let's say 1Gb). How would we go about building such a tool? What are some questions we should ask the client? How would we provide the required functionality? Is it even possible? (3)

END OF QUIZ