

START OF QUIZ

Student ID:

74035403,Zheng,Meiyu

Question 1

Topic: Lecture 8

Source: Lecture 8

Imagine that we have a dialogue system trained with reinforcement learning. What part of a dialogue might result in a negative reward (ie, a penalty) to the system's policy algorithm?
(2)

Question 2

Topic: Lecture 6

Source: Lecture 6

Explain the purpose of mean reciprocal rank, and how it works. (1)

Question 3

Topic: Lecture 5

Source: Lecture 5

Briefly describe a “factoid-based” question, and one way that a QA system might answer it.
(1)

Question 4

Topic: Lecture 7

Source: Lecture 7

Generate a frame for a "recommend a movie" dialogue action. It should have at least 5 slots to fill. (2)

Question 5

Topic: Lecture 5

Source: Lecture 5

Explain why a good IR tool is necessary to perform Bert-based Q/A. (1)

Question 6

Topic: Lecture 8

Source: Lecture 8

We waited until the last week of classes to talk about policy-making systems (like the one in ChatGPT), but several other systems you've looked at over the program could be considered to have a policy algorithm in place. Briefly describe one, and how you view it as a decision policy. (2)

Question 7

Topic: Lecture 7

Source: Lecture 7

We discussed slot error rate in class, but it's fully-supervised. Can you think of a distantly-supervised way to calculate essentially the same thing? (1)

Question 8

Topic: Lecture 6

Source: Lecture 6

Provide a reasonable logical representation of the question “Who starred in Casablanca?” (1)

Question 9

Topic: Coding

Source: Lecture 5

Imagine that we are using a Q/A system for movie recommendation (by asking questions like “What is a good movie like Shawshank Redemption?”). Bert is likely not going to be sufficient to answer this question. Describe how you could modify the Bert Q/A reader to find good answers. (3)

END OF QUIZ