

START OF QUIZ

Student ID:

88068333, Wang, Junrui

Question 1

Topic: Lecture 5

Source: Lecture 5

We've seen co-occurrence matrices weighted by TF-IDF- would it make sense to weight them by PMI? Briefly explain. (1)

Question 2

Topic: Lecture 8

Source: Lecture 8

Do you think we could use word embeddings for coreference resolution? What kind of assumptions would we be making, and why do you think it might still be a very difficult task?
(2)

Question 3

Topic: Lecture 8

Source: Lecture 8

What is the purpose of an antecedent in anaphoric resolution? (1)

Question 4

Topic: Lecture 6

Source: Lecture 6

Why do we say that the analogy task is an "intrinsic" evaluation of our word embeddings? (1)

Question 5

Topic: Lecture 5

Source: Lecture 5

In class, we talked about how a "typical" dimensionality for embeddings is in the range of 100-500. What might be some consequences if we estimated too low or too high? (2)

Question 6

Topic: Lecture 6

Source: Lecture 6

What is the purpose of negative sampling in a Word2Vec model? (1)

Question 7

Topic: Lecture 7

Source: Lecture 7

The TextTiling algorithm we looked at just looked at raw word overlap (possibly with stop-word removal). Describe a way that we could improve the algorithm to maximize coverage.
(2)

Question 8

Topic: Lecture 7

Source: Lecture 7

Explain the underlying assumption of the TextTiling algorithm. (1)

Question 9

Topic: Coding

Source: Coding

Write a short function that tries to find a good value for k in truncated SVD. You'll essentially be writing your own version of the evaluate word analogies function. For each of our the analogies, you'll need to do the vector math we were doing in the capital city determination, and return the 1-closest vector - if it's what we're looking for, it's correct. If not, it's wrong. (3)

END OF QUIZ