

START OF QUIZ

Student ID:

30542179, Yang, Yimei

Question 1

Topic: Lecture 7

Source: Lecture 7

How might you modify a standard sentiment analyzer to track change in sentiment over time? (2)

Question 2

Topic: Lecture 5

Source: Lecture 5

SVM ranking takes advantage of the fact that an ordinal problem can be transformed into a binary "larger than" problem by simple subtraction of feature vectors. It's typically done with a linear SVM. Do you think we could apply a similar trick with a neural model? Why or why not? (2)

Question 3

Topic: Lecture 6

Source: Lecture 6

We saw that age and gender are relatively easy to predict from tweet history, but that personality traits are a lot harder. Why do you think that is? (1)

Question 4

Topic: Lecture 8

Source: Lecture 8

What properties of code-switched text are useful for identifying the language of the text?
(List at least 2) (1)

Question 5

Topic: Lecture 8

Source: Lecture 8

What is one similarity and one dissimilarity between emojis and emoticons? (1)

Question 6

Topic: Lecture 5

Source: Lecture 5

In class, we said that "fake" fake reviews are often too prototypical when they are generated by hand. Given the tools you're familiar with, how do you think we could generate fake reviews automatically? Do you think they would suffer from the same problem? (2)

Question 7

Topic: Lecture 7

Source: Lecture 7

Why is datetime functionality necessary? That is, why can't we just use the date and time separately? (1)

Question 8

Topic: Lecture 6

Source: Lecture 6

What is distant supervision, and why can we apply it to social media? (1)

Question 9

Topic: Long

Source: Lecture 7

In class, we talked about the granularity of choropleths, and how a world map might not capture intricacies of a smaller group, simply lumping them in with other people. How would you balance the tradeoff between finer granularity and the need to collect more individualized (and potentially private) data? (3)

END OF QUIZ