

START OF QUIZ

Student ID: 29404548, Lee,
Ai Ka

Question 1

Topic: Lecture 8

Source: Lecture 8

How does ICL differ from fine-tuning? (1)

Question 2

Topic: Lecture 6

Source: Lecture 6

If you were building your own parallel corpus, what kind of information would you prioritize?
What questions would you ask your stakeholders? (1)

Question 3

Topic: Lecture 8

Source: Lecture 8

Explain the role of clustering when performing self-training? (1)

Question 4

Topic: Lecture 7

Source: Lecture 7

Describe situations where margin based selection is more and less appropriate, and why. (1)

Question 5

Topic: Lecture 6

Source: Lecture 6

From your perspective, what is the biggest advantage and disadvantage of open-source models? (1)

Question 6

Topic: Lecture 7

Source: Lecture 7

You've been using AL with multiple annotators. As a sanity check, you have several instances labeled by multiple annotators, but find that the annotations are inconsistent. How can you remedy the problem and select good examples, without knowing the language you are having annotated? (2)

Question 7

Topic: Lecture 5

Source: Lecture 5

Imagine we have a multilingual encoder-model like mBERT, and a multilingual decoder-only model. Do you think we could train the encoder on one set of languages, and then the decoder on a larger set, and better understand the new languages? What kind of adaptations would need to be done? Do you think it would improve zero-shot learning on languages not included in either? (2)

Question 8

Topic: Lecture 5

Source: Lecture 5

You're working with MT5, and you find it's not doing very well on your target language, even after fine-tuning. What do you do? Would your answer change if the model were mBert, instead? (2)

Question 9

Topic: Long

Source: Lecture 5

Imagine you're working on adapting a multilingual LLM for a government that wants it to operate fluently in 10 national languages, including both high- and low-resource languages, and avoid colonial-language bias. Describe a fine-tuning and evaluation pipeline that could help adapt the model fairly across languages. What ethical and linguistic challenges might arise, and how would you mitigate them? How would you include community feedback in the loop? (3)

END OF QUIZ