START OF QUIZ
Student ID:
31181845,Yuan,Su

Question 1

Topic: Lecture 8
Source: Lecture 8

How does ICL differ from fine-tuning? (1)

Question 2

Topic: Lecture 8
Source: Lecture 8

How does silver data differ from synthetic data? (1)

Question 3

Topic: Lecture 6
Source: Lecture 6

If you were building your own parallel corpus, what kind of information would you prioritize? What questions would you ask your stakeholders? (1)

Question 4

Topic: Lecture 7
Source: Lecture 7

Why is entropy a good measure to use when using QbC? (1)

Question 5

Topic: Lecture 5
Source: Lecture 5

Why do MLLMs tend to eventually see a decrease in quality on HRLs? (1)

Question 6

Topic: Lecture 7
Source: Lecture 7

We discussed active learning with respect to classification, but what about regression tasks? What similarities / differences might make active learning suitable or unsuitable to regression? (2)

Question 7

Topic: Lecture 5
Source: Lecture 5

You're working with MT5, and you find it's not doing very well on your target language, even after fine-tuning. What do you do? Would your answer change if the model were mBert, instead? (2)

Question 8

Topic: Lecture 6
Source: Lecture 6

In the lab, you likely saw that a certain part took much longer than others, and produced noisy output. How might you adjust your strategy, speed up the methodology, and what simple tools could you use to decrease noise? (2)

Question 9

Topic: Long
Source: Lecture 5

Imagine you're working on adapting a multilingual LLM for a government that wants it to operate fluently in 10 national languages, including both high- and low-resource languages, and avoid colonial-language bias. Describe a fine-tuning and evaluation pipeline that could help adapt the model fairly across languages. What ethical and linguistic challenges might arise, and how would you mitigate them? How would you include community feedback in the loop? (3)

END OF QUIZ