

START OF QUIZ

Student ID:

88068333, Wang, Junrui

Question 1

Topic: Lecture 8

Source: Lecture 8

Why don't we use a higher-order language model to perform IR? (1)

Question 2

Topic: Lecture 6

Source: Lecture 6

Imagine we performed LDA on the classes in this block. What might their Beta distributions look like? (2)

Question 3

Topic: Lecture 8

Source: Lecture 8

What do we mean by interpolation? (1)

Question 4

Topic: Lecture 5

Source: Lecture 5

Why do we need methods like t-SNE? (1)

Question 5

Topic: Lecture 7

Source: Lecture 7

Why do we generally care more about precision than recall in IR? (1)

Question 6

Topic: Lecture 5

Source: Lecture 5

We often weight our matrices using something like PMI or TF-IDF. Do you think it would make sense to do this after applying SVD? Why or why not? (2)

Question 7

Topic: Lecture 6

Source: Lecture 6

Why can't we just run an HMM over documents to discover the latent states like we do for POS-tagging? (1)

Question 8

Topic: Lecture 7

Source: Lecture 7

Explain why the cosine similarity between a document and query vector is roughly equivalent to adding up the TF-IDF scores of each word in the document that occurs in the query.
(2)

Question 9

Topic: Coding

Source: Coding

Write a function that returns the most likely n documents given a term-document matrix, a smoothing parameter, and a query. (3)

END OF QUIZ